

Per-module inference latency on Apple M4 Max

